

# Time Series Analysis

Using an Event Streaming Platform

***Virtual Meetup - September 8-th 2020***

Dr. Mirko Kämpf - Solution Architect @ *Confluent, Inc.* - Team CEMEA

# ***Abstract***

## **Time Series Analysis using an Event Streaming Platform**

Advanced time series analysis (TSA) requires very special data preparation procedures to convert raw data into useful and compatible formats.

In this presentation you will see some typical processing patterns for time series based research, from simple statistics to reconstruction of correlation networks.

The first case is relevant for anomaly detection and to protect safety.

Reconstruction of graphs from time series data is a very useful technique to better understand complex systems like supply chains, material flows in factories, information flows within organizations, and especially in medical research.

With this motivation we will look at typical data aggregation patterns. We investigate how to apply analysis algorithms in the cloud. Finally we discuss a simple reference architecture for TSA on top of the Confluent Platform or Confluent cloud.

This presentation is about linking:

- *Time-Series-Analysis (TSA)*
- *Network- or Graph-Analysis*
- *Complex Event Processing (CEP).*

*Confluent Platform*

Research work ends often with nice charts, scientific papers, and conference talks.

But, many published results **can't** be reproduced -  
often because the setup it is simply too complicated ...

**Question:**

How can we integrate data streams,  
experiments, and decision making better?

# Why not using batch processing?

Study anatomy ...



- **Batch processing is fine:**
  - as long as your data doesn't change.
  - in PoCs for method development in a lab.
  - for research in a fixed scope.

# Why using Kafka?

- **Stream processing is better:**

- for real time business in changing environments.
- iterative (research) projects.
- repeatable experiments on replayed data.

Study and influence  
the living system ...



# Content:

## **(1) Intro**

Typical types of event

How to identify hidden events?

## **(2) The Challenge**

## **(3) Approach**

Time Series Analytics &  
Network Analytics in Kafka

Create time series from events

Create graphs from time series pairs

## **(4) Demo**

**(5) Architecture & Building Blocks:**

**(6) Unified Domain Model**

# Events - 1

## Business events

- transaction records
- discrete observation

## How to handle events?


JUTS SIMPLE  
OBSERVATION &  
DATA CAPTURING



**A Sale**



**An Invoice**



**A Trade**



**A Customer  
Experience**


# Events - 2

Well defined events

- in known context

## How to identify events?

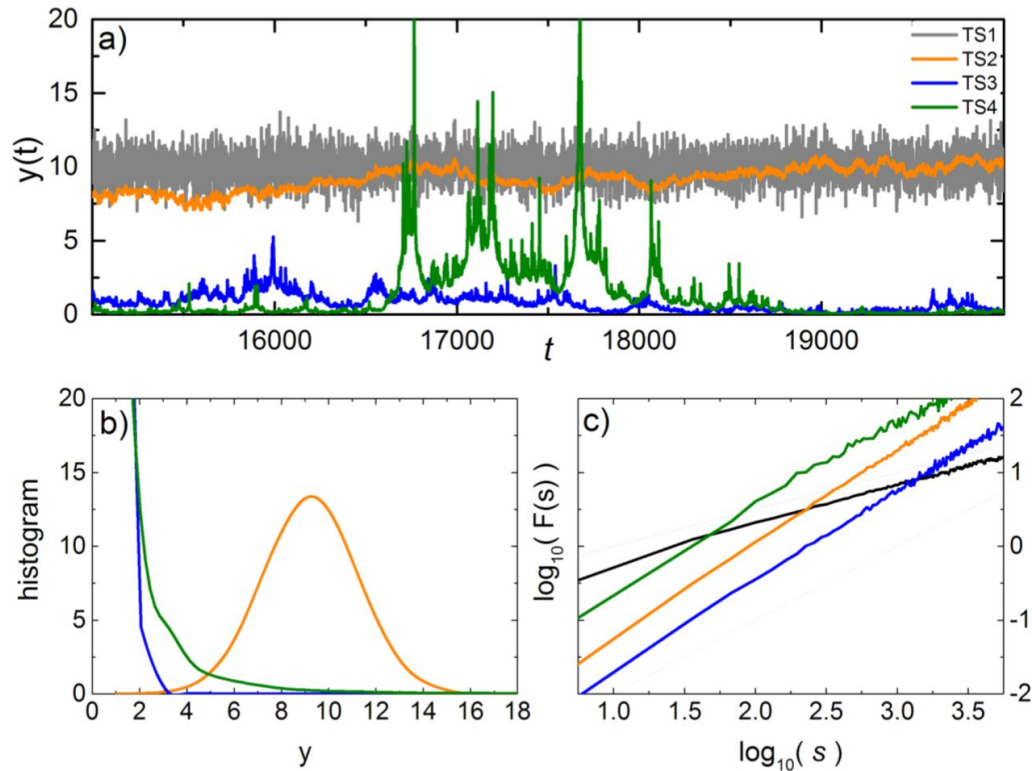
Sometimes: SIMPLE  
Sometimes: DATA ANALYSIS





# Univariate TSA: single episodes are processed

- Extreme values
- Patterns
- Distribution of values
- Fluctuation properties
- Long-term correlations (memory effects)



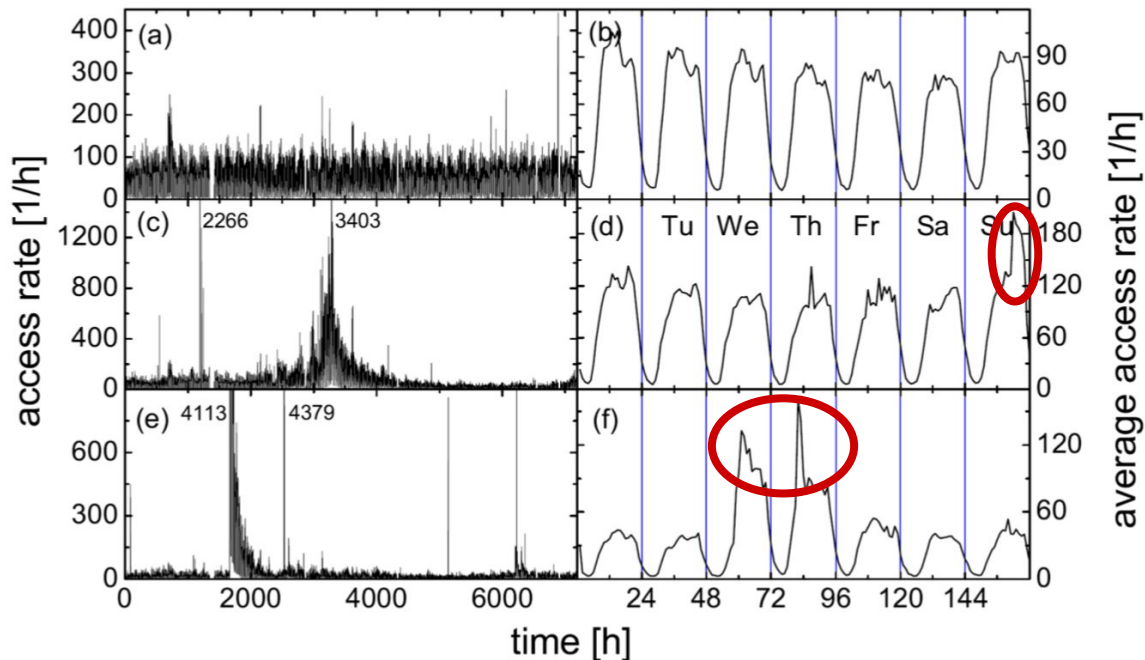
# Events - 3

## Extreme Events

- “outliers” in unknown context


## How to handle?

ADVANCED  
DATA ANALYSIS (& ML)



# Multivariate TSA: pairs / tuples of episodes are processed

- Comparison Similarity measures for link creation



Distribution of cross-correlation coefficients for pairs of access-rate time series of Wikipedia pages (top) compared to surrogate data (bottom) - 100 shuffled configurations are considered

# Reality is Complex:

**We should simplify a bit!**


Simplification in our method can lead to isolation:

- DATA SILOS
- OPERATIONAL SILOS


**SOLUTION:**

GRAPHS capture structure.


TIME SERIES capture properties over time (history).

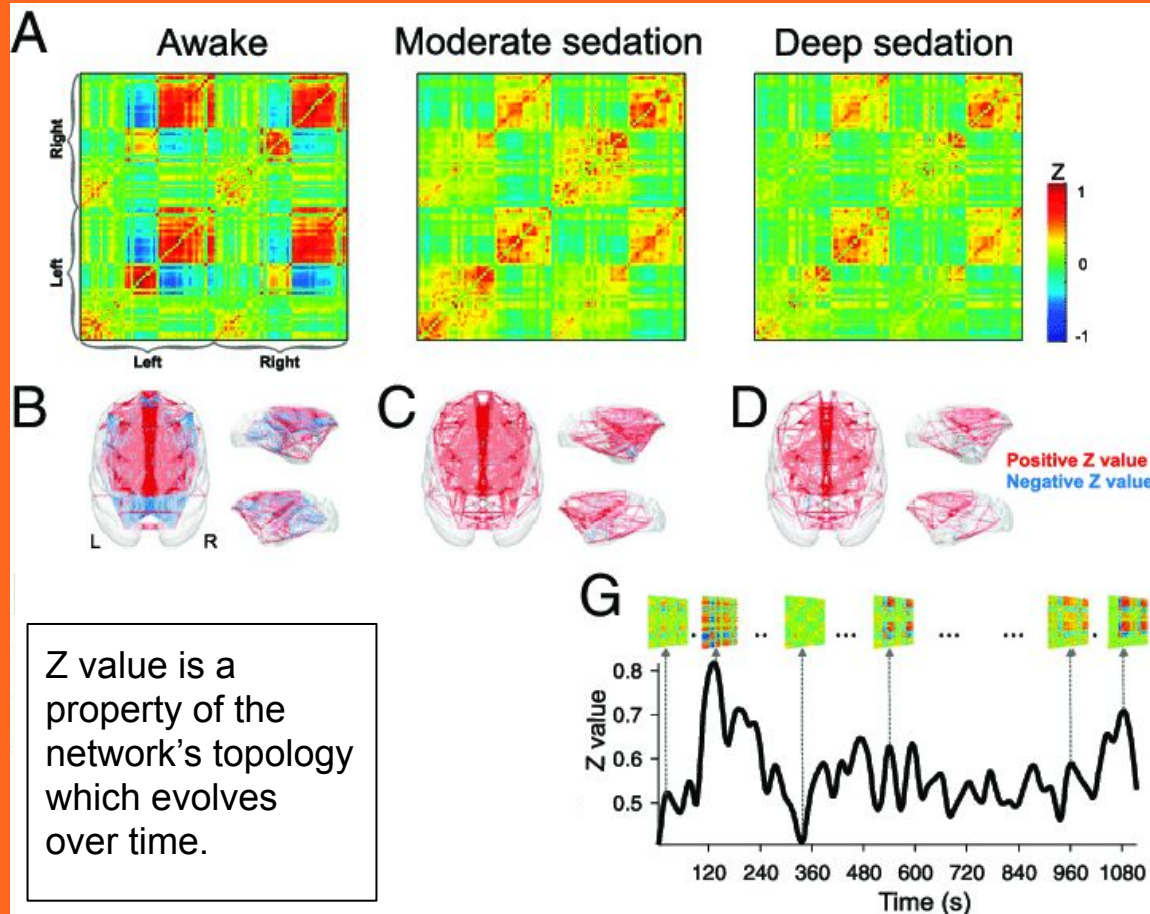


relations as graph or matrix:



objects in groups:





WHAT?

WHY?

# Recap:

What events are  
and how to process  
event-data can be  
misunderstood  
*or simply unclear.*

**It all depends on our view  
and our goals!**

## IT Operations

- Server crash
- Cyber crime

Special procedures  
are established.

## Business Events

- Big deal won
- Technical issue solved

The events which make  
people & the market  
happy :-)

## Transactions (in business)

- orders placed
- products shipped
- bills paid

Event Driven Architecture

## Extreme Events:

- Service slow down due to emerging bottlenecks
- Increased demand in a resource

Complex Event Analysis

# The Challenge:

How can you combine unbound data assets and scientific methods?

- You bring data to a place where it can be processed easily, for example to a cloud system or into special purpose systems, such a event processing platform.
- You integrate important algorithms as early as possible in your processing pipeline to get results fast, using streaming processing capabilities.

—

**Things can become complicated:**

- Complex Event Analysis**
- Integration Across Domains**
- Extraction of Hidden Event**

—





## METHODOLOGY

A short, thick white horizontal bar.

# Complex Event Analysis

- time series analysis and ML reveal hidden events
  - multi-stage processing is usually needed
- 
- A short, thick white horizontal bar.

## ORGANIZATION & OPERATIONS

---

# Integration Across Domains

- distributed event processing systems are used
  - apps consume and produce events of different flavors
  - event-types and data structures may change over time
-

# Problems on ORGANIZATION level:

Many legacy systems can't be integrated without additional expensive servers. Often, this data is unreachable for externals.

Business data is managed by different teams using different technologies.

Data scientists use data in the cloud, and they all do really L❤️VE notebooks. But often, they don't use any automation.

## TECHNOLOGY & SCIENCE

---

# Extraction of Hidden Events

- requires **Applied Data Analysis & Data Science**
  - embedding of **Complex Algorithms** in IT landscape
  - integration of **GPU/HPC** and streaming data pipelines
-

# Kafka and its Ecosystem ...


- are considered to be middleware, managed by IT people:
  - researchers do not plan nor build their experiments around this technology yet.
- don't offer ML / AI components:
  - many people think, that a model has to be executed on an edge-device or in the cloud.

*Just  
because  
they don't  
understand  
doesn't mean  
you're on the  
wrong path.*

@QWORLDSTAR

# Yes, Apache Kafka can support agile (data)experiments!

- **Kafka APIs** give access to data (flows) in real time.
  - allows replay of experiments at a later point in time
- Kafka allows **variation of analysis without redoing a simulation or ingestion** by simply reusing persisted event-streams again.
- **Kafka Streams** and **KSQL** allow data processing in place:
  - this allows faster iterations, for example because plausibility checks can be done in place
  - the streaming API gives freedom for extension of core logic
  - DSL and KSQL will save you a lot of implementation time





Why not building something great using the right tools ???

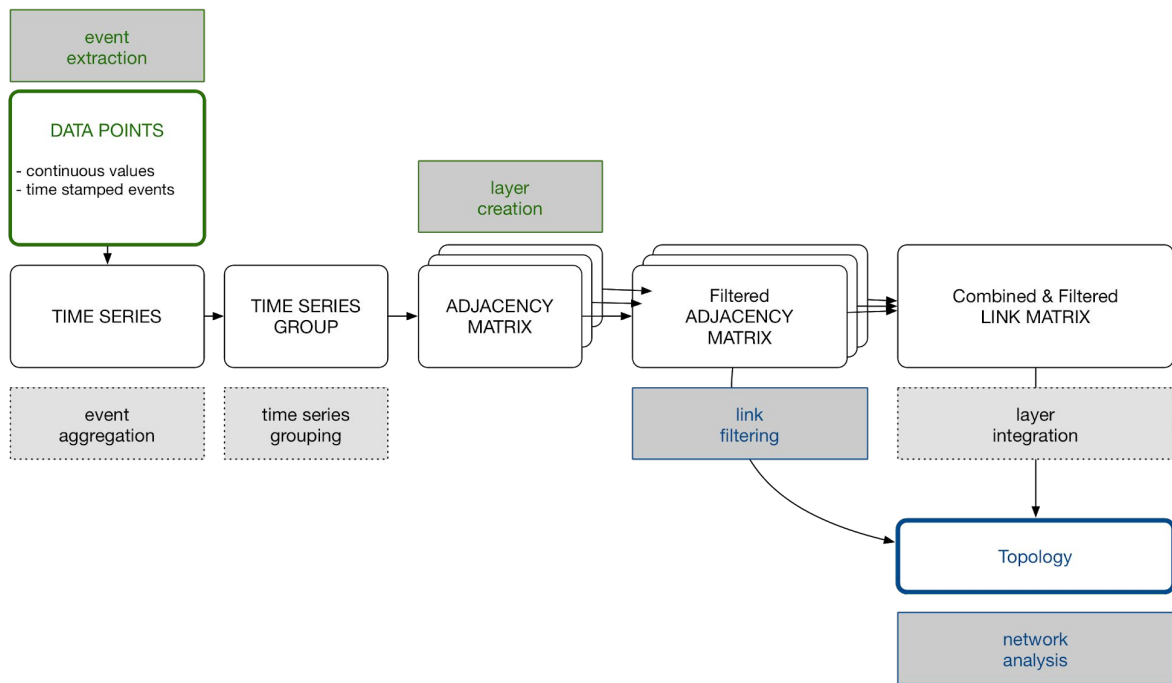
# ADVANCED TIME SERIES ANALYSIS & NETWORK ANALYSIS

... how does it work?



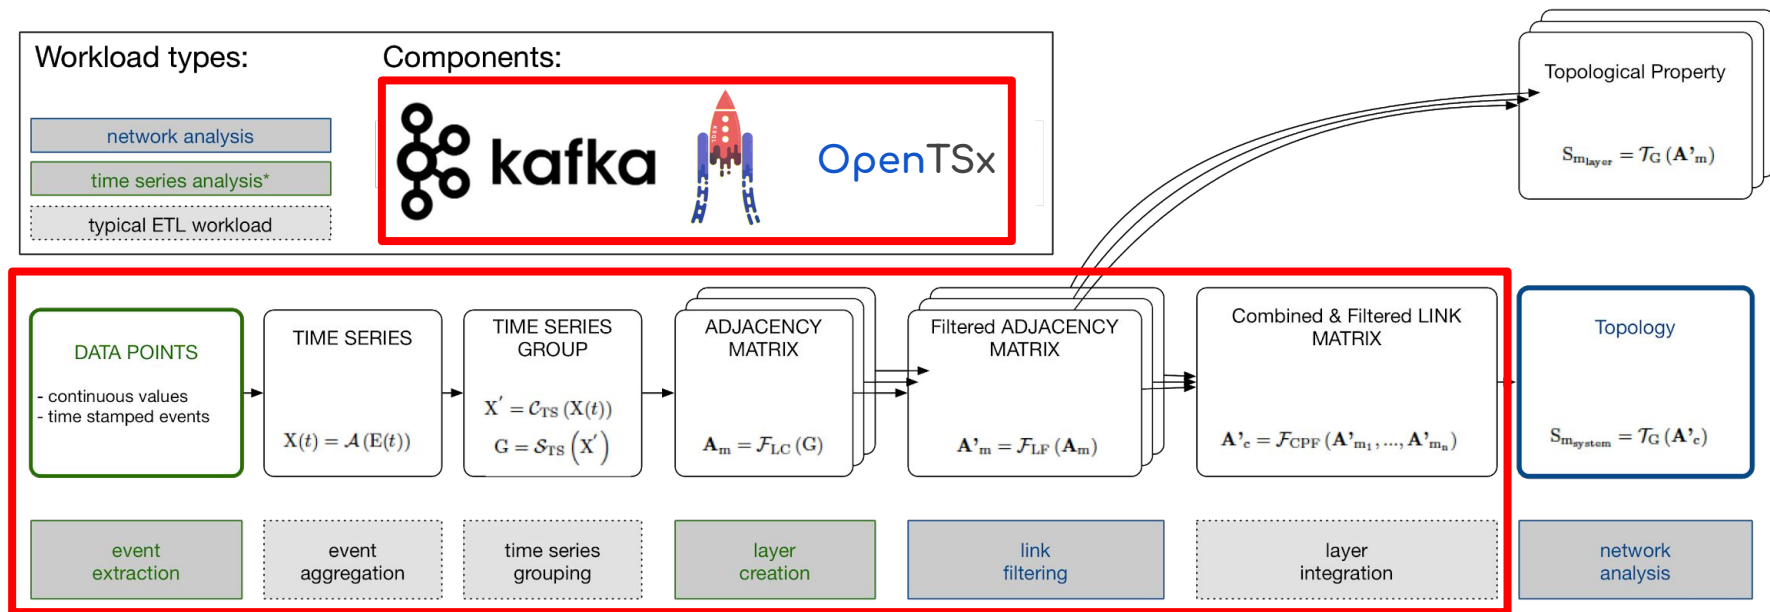
# Network Reconstruction & Topology Analysis:

## The Approach



# Network Reconstruction & Topology Analysis:

## A Standardized Event Processing Pipeline




# ADVANCED TIME SERIES ANALYSIS & NETWORK ANALYSIS

... and how does this fit into Kafka?

**The following slides will show  
how TSA concepts  
and Kafka concepts  
fit together.**


# Table Stream Duality



BIRTE '18, August 27, 2018, Rio de Janeiro, Brazil

M.J. Sax, G. Wang, M. Weidlich, J.-C. Freytag

# Create Time Series from Event Streams: By Aggregation, Grouping, and Sorting



# From Table of Events to - Time Series

**Table 1: Operators with their input and output types**

Operator	1st Input	2nd Input	Output
filter, mapValue	KStream		KStream
	KTable		KTable
map, flatMap	KStream		KStream
groupBy → agg	KStream		KTable
	KTable		KTable
groupBy + windowBy → agg	KStream		KTable
inner-/left-/outer-join	KStream	KStream	KStream
inner-/left-/outer-join	KTable	KTable	KTable


## Table Stream Duality $\Rightarrow$ Time Series and Graphs

A ***time series*** is a table of **ordered observations** in a fixed context.

A ***graph*** can be seen as a table of node- and link-properties - stored in **two tables**.

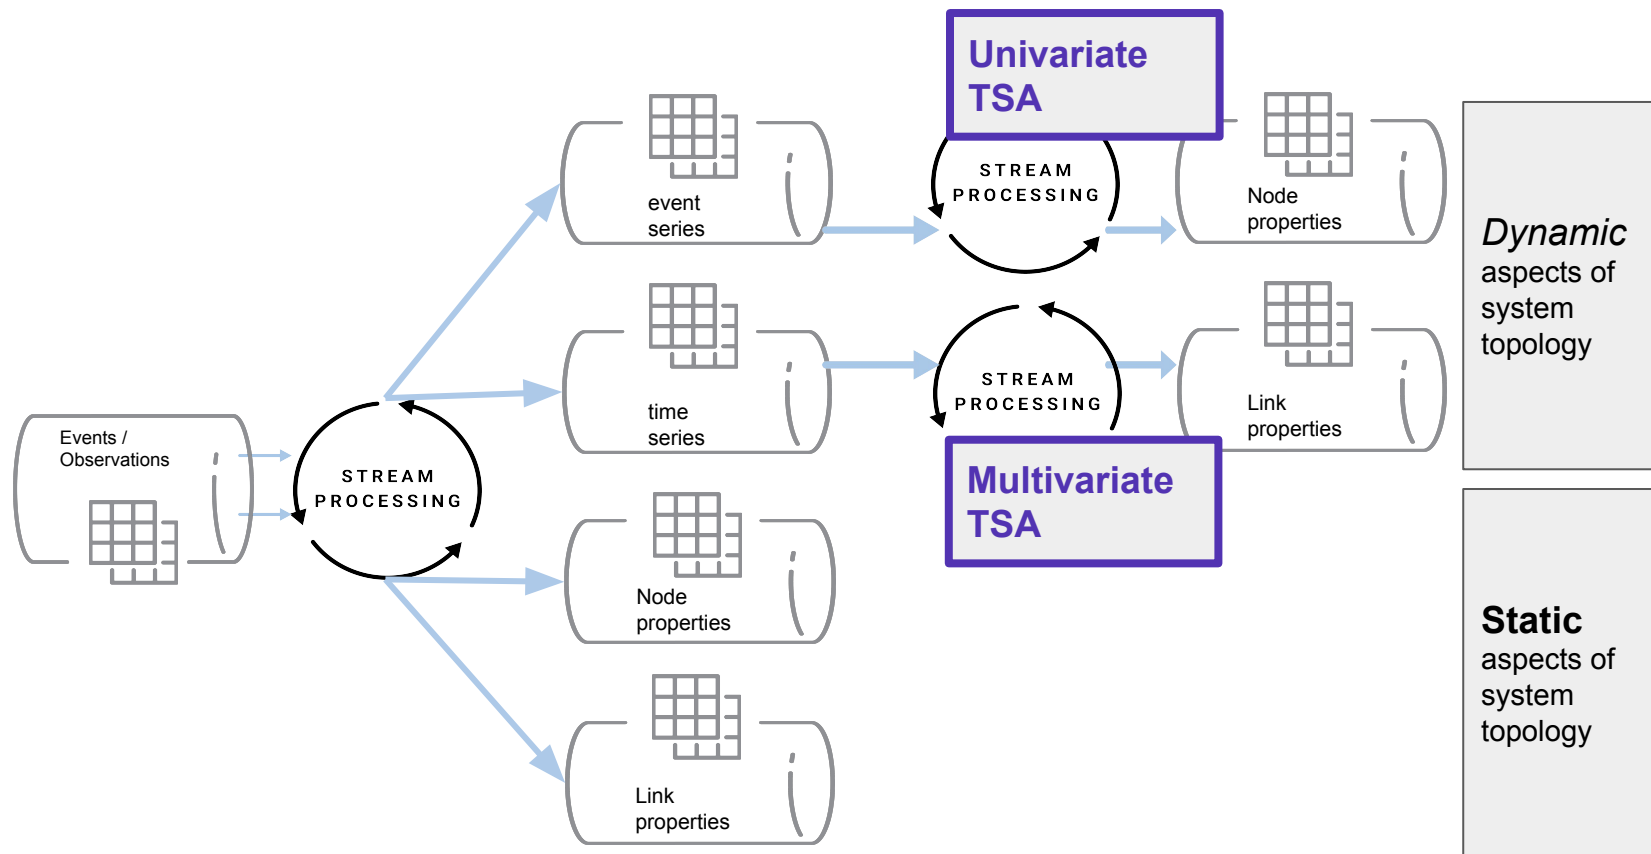


# Create Networks from Event Streams: By Aggregation, Grouping, and Sorting

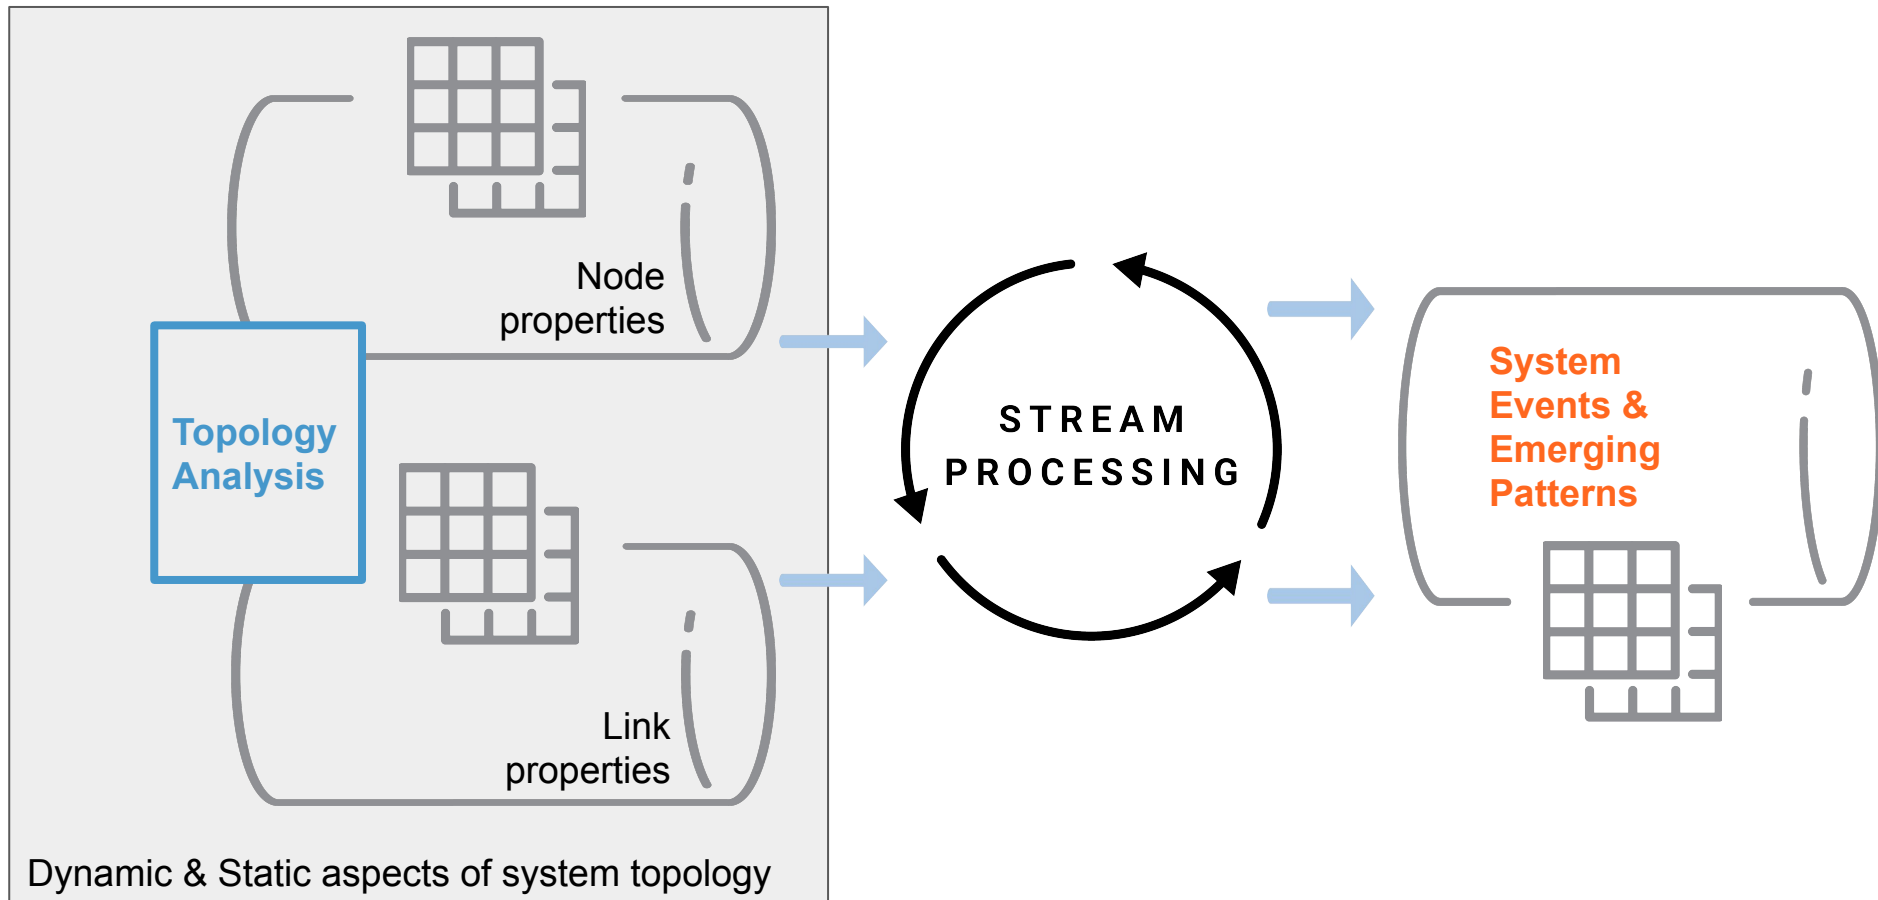


# Multi Layer Stream Processing:

## TSA to Reveal Hidden System Properties



# Complex Event Processing: For Complex Systems



# Use the Table-Network Analogy: **Kafka Graphs**

## Creating Graphs

**large durable graphs:**

Persisted in Kafka topic

A graph in Kafka Graphs is represented by two tables from Kafka Streams, one for vertices and one for edges. The vertex table is comprised of an ID and a vertex value, while the edge table is comprised of a source ID, target ID, and edge value.

```
KTable<Long, Long> vertices = ...  
KTable<Edge<Long>, Long> edges = ...  
KGraph<Long, Long, Long> graph = new KGraph<>(  
    vertices,  
    edges,  
    GraphSerialized.with(Serdes.Long(), Serdes.Long(), Serdes.Long())  
);
```


<https://github.com/rayokota/kafka-graphs>

# Sliding Windows: Define the Temporal Graphs

In some use cases, we don't want to keep the node and link data in topics:

- nodes aren't always linked
- changes are very fast
- focus on activation patterns, rather than on underlying structure

It is fine to calculate the correlation links and the topology metrics on the fly, just for a given time window.



# Back to Streams



# Demo ...



```

public static void main(String[] args) throws Exception {

    md = MessageDigest.getInstance("MD5");

    StreamsConfig streamsConfig = new StreamsConfig( props );
    StreamsBuilder builder = new StreamsBuilder();

    Serde<String> stringSerde1 = Serdes.String();
    Serde<String> stringSerde2 = Serdes.String();
    final Map<String, String> serdeConfig = Collections.singletonMap("schema.registry.url",
        props.getProperty( "schema.registry.url" ));

    final Serde<Event> valueSpecificAvroSerde = new SpecificAvroSerde<>();
    valueSpecificAvroSerde.configure(serdeConfig, isKey: false); // `false` for record values

    KStream<String, Event> my_events = builder.stream( topic: "OpenTSx_Events", Consumed.with( stringSerde1, valueSpecificAvroSerde));

    my_events = my_events.mapValues( (v) -> getCompactedEvent(v) );
    my_events.to( topic: "EVENT_DATA_TOPIC", EXPERIMENT_TAG, Produced.with( stringSerde2, valueSpecificAvroSerde ) );
    my_events.print(Printed.<>toSysOut().withLabel("[EVENT_DATA]"));

    Topology topology = builder.build();
    dumpTopology(topology);

    KafkaStreams kafkaStreams = new KafkaStreams(topology, streamsConfig);
    kafkaStreams.setUncaughtExceptionHandler((t, e) -> { LOG.error("had exception ", e); });

    LOG.info("Starting event processing example { experiment_duration=" + experiment_duration + " ms}.");
    kafkaStreams.cleanUp();
    kafkaStreams.start();

    if ( experiment_duration > 0 ) { Thread.sleep(experiment_duration); }

    LOG.info("Shutting down the event processing example application now.");
    kafkaStreams.close();

}

private static void dumpTopology(Topology topology) throws IOException {
    FileWriter fw = new FileWriter( fileName: props.getProperty(StreamsConfig.APPLICATION_ID_CONFIG) + "_topology.dat");
    fw.write( topology.describe().toString() );
    fw.flush();
    fw.close();
}

```

The project will provide reusable streaming apps so that developers implement just their specific logic.



```
package org.opentsx.udf;

import io.confluent.ksql.function.udf.Udf;
import io.confluent.ksql.function.udf.UdfDescription;
import io.confluent.ksql.function.udf.UdfParameter;

/**
 * select mdExtract( TSTART, TEND ) as LENGTH FROM EPISODES_STREAM EMIT CHANGES;
 */
@UdfDescription(name      = "mdExtract",
               description = "Extracts metadata from single episods.",
               author      = "Mirko Kämpf",
               version     = "3.0.1" )

public class EpisodesMetadataExtractor {

    @Udf(description = "Calculates length of an episode (in ms).")
    public double calc_length(
        @UdfParameter(value = "TSTART", description = "START time") final long TSTART,
        @UdfParameter(value = "TEND", description = "END time") final long TEND
    )
    {
        return (double)(TEND - TSTART);
    }
}
```

```
/**  
 * Simple threshold filter  
 *  
 * @param k – message key (is ignored in this example)  
 * @param e – message value (our event)  
 *  
 * @return boolean value for filtering  
 */  
private static boolean getCriticalEvents(String k, Event e) {  
    if ( e.getValue() > 3.0 ) return true;  
    else return false;  
}
```

```

/**
 * getSimpleStats from TS0
 * -----
 *
 * @param k - key (we extract the uuid which is part of the key and put it also into our result map.
 * @param v - the result map, calculated by the OpenTSx library. Simple descriptive statistics in this case.
 *
 * @return Hashtable with results encoded as JSON.
 */
private static String getSimpleStats(String k, EpisodesRecord v) {

    TimeSeriesObject tso = TimeSeriesObject.getFromEpisode( v );

    Hashtable<String,String> stats = tso.getStatisticData( new Hashtable<String,String>() );

    Map<String, String> retMap = new Gson().fromJson(
        k, new TypeToken<HashMap<String, String>>() {}.getType()
    );

    stats.put( "mkey" , retMap.get("uuid") );

    return gson.toJson( stats );

}

```

**Finally, we can compose a complex flow ...**

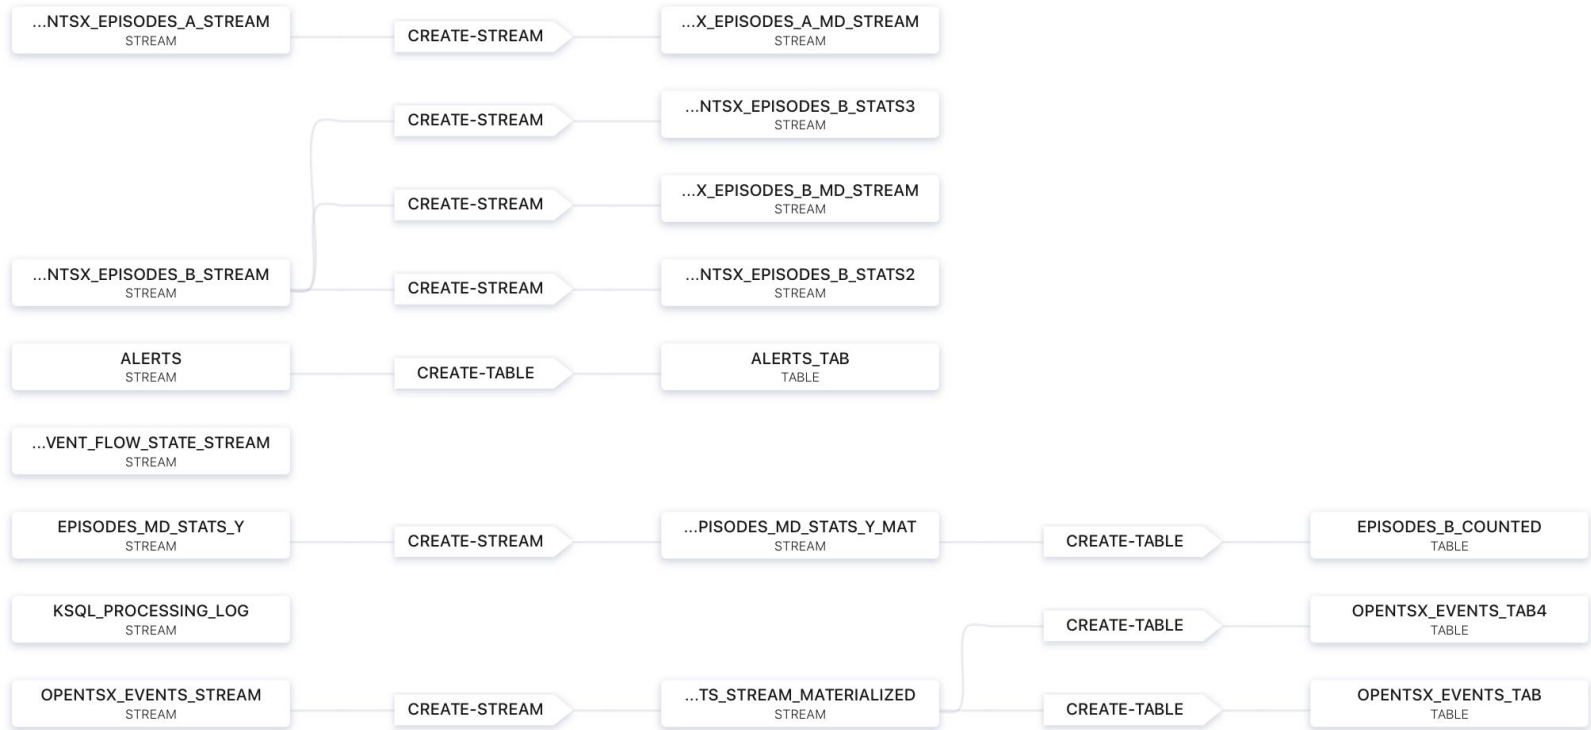
# ksqlDB

Editor   **Flow**   Streams   Tables   Running queries

 Search

```
ksql> select * from alerts_tab where URI='http://www.w3.org/TR/2004/metrics-owl-20200210/deviceID#0815a';
```

URI	WINDOWSTART	WINDOWEND	Z
http://www.w3.org/TR/2004/metrics-owl-20200210/deviceID #0815a	1599496830000	1599496845000	234
http://www.w3.org/TR/2004/metrics-owl-20200210/deviceID #0815a	1599496845000	1599496860000	9677
http://www.w3.org/TR/2004/metrics-owl-20200210/deviceID #0815a	1599496860000	1599496875000	11705



# Demo: OpenTSx

Generate some events and some episodes.

Apply some time series processing procedures on:

- a stream of episodes.
- a stream of events


Complex procedures are composed from a set of fundamental building blocks.

Visualize the flows and dependencies.

# Let's Remember ...

# We use 4 building blocks ...

Domain specific logic is implemented in small reusable components:




*Domain Driven Design*

Source Connectors integrate sources ...



*Legacy and Future Systems*

Data flows are no longer transient.  
The event log acts as single source of truth.



*Paradigm Shift in Data Management*

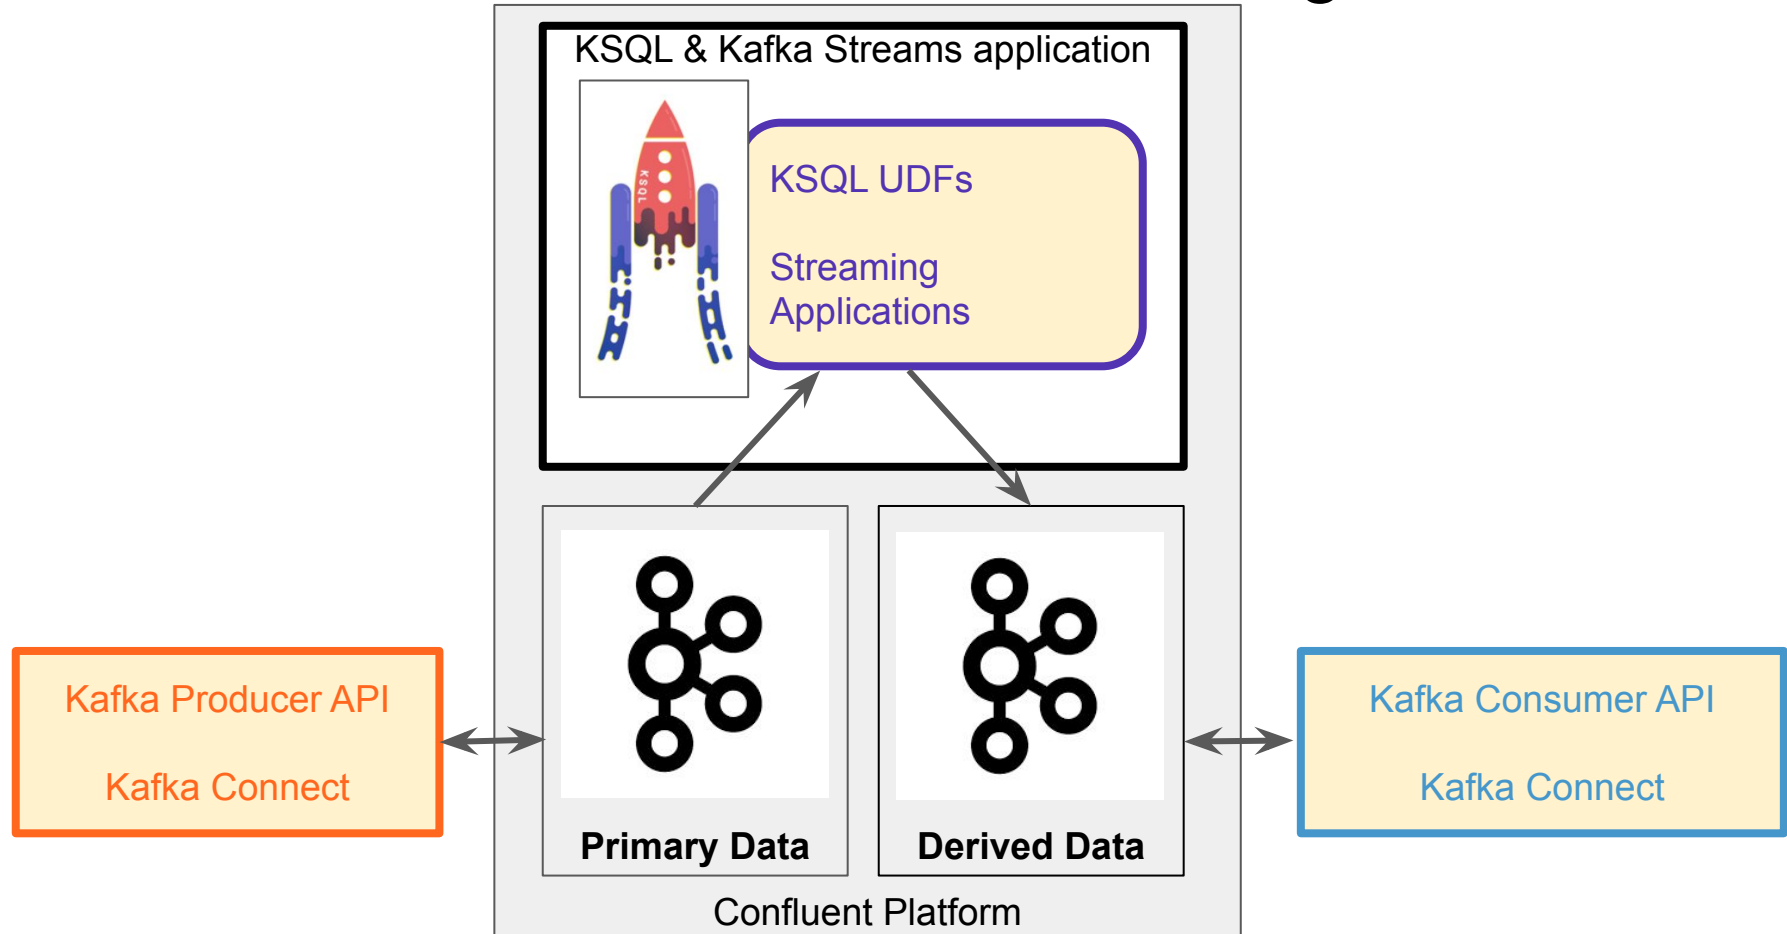
Sink Connectors integrate external targets ...



*Special Purpose Systems*



# ... for our Time Series Processing Platform



# Summary:

Because Kafka is a scalable & extensible platform it fits well for complex event processing in any industry on premise and in the cloud.

Kafka ecosystem provides extension points for any kind of domain specific or custom functionality - from advanced analytics to real time data enrichment.

Complex solutions are composed from a few fundamental building blocks:

# What to do next?

- (A) Identify **relevant main flows** and **processing patterns** in your project.
- (B) Identify or implement source / sink **connectors** and **establish 1st flow**.
- (C) Implement **custom transformations** as **Kafka independent components**.
- (D) Integrate the processing topology as Kafka Streams application:
  - (a) Do you apply standard transformations and joins (for enrichment)?
  - (b) Is a special treatment required (advanced analysis)?
  - (c) Do you need special hardware / external services (AI/ML for classification)?
- (E) Share your connectors and UDFs with the growing Kafka community.
- (F) Iterate, add more flows and more topologies to your environment.

**END**

**Thank you !**

**mirko@confluent.io**



**@semanpix**